

# Privacy-Preserving Functional Database Exploration with Lattigo

Jean-Philippe Bossuat

jeanphilippe.bossuat@gmail.com



# LATTIGO

*Disclaimer: The views and opinions expressed in this talk are my own and they do not necessarily represent the views and opinions of my employer.*

# Table of Contents

1 INTRODUCTION

2 THEORY & IMPLEMENTATION

3 PERFORMANCE

# INTRODUCTION



- This presentation is aimed at showcasing how LATTIGO can be used to implement a very diverse circuit solving a practical use case
- The constructions used in this presentation are based on a joint work with **Malika Izabachène**, that will be accessible on ePrint in the coming days
- LATTIGO is an open-source Go module that implements RLWE based homomorphic-encryption primitives and Multiparty-Homomorphic-Encryption (MHE) based protocols
- LATTIGO was originally developed at the EPFL Laboratory for Data Security by **Jean-Philippe Bossuat** and **Christian Mouchet**, supervised by **Juan Troncoso**, and since 2022 it has been maintained by **Tune Insight SA**

- A scientist (the Client) would like to conduct a large-scale medical study requiring patients with **specific combinations of attributes**
- To be funded, the scientist must first conduct a preliminary feasibility study to **assess if there are enough subjects available** in a patient database (the Server) meeting the study criteria
- To protect the **IP of the study** the Server should not learn anything about the selection criteria
- To protect the **privacy of the patients**, the Client should only learn a binary value: if there are enough patients meeting the selection criteria in the Server's database.

Given:

- A database  $\mathcal{P}$  that consists in a  $n \times h$  patient matrix
- A list of  $h$  scoring functions
- An individual score threshold  $t_0$
- A global count threshold  $t_1$

We want to answer the following question in a privacy preserving (and efficient) way:

*Does  $\mathcal{P}$  contains at least  $t_1$  patients  
whose individual score is at least equal to  $t_0$ ?*

# THEORY & IMPLEMENTATION

## Overview of the circuit:

$$\text{Thresh}_{\text{Enc}(t_1)} \left( \sum_{i=0}^{n-1} \text{Thresh}_{\text{Enc}(t_0)} \left( \sum_{j=0}^{h-1} \text{Enc}(f_j) \otimes \mathcal{P}[i][j] \right) \right)$$

- $\mathcal{P}$ :  $n \times h$  database
- $f_j$ : attribute scoring functions  $\mathbb{R}_{[a,b]} \rightarrow \mathbb{Z}_p$
- $t_0$ : individual selection threshold
- $t_1$ : global selection threshold
- $\text{Tresh}_t(x) = 1$  if  $x \geq t$  else 0





This circuit is interesting because it makes use of:

- **Encrypted Lookup Tables:** evaluation of the private scoring functions  $\text{Enc}(f_j)$
- **Ring-Packing:** repacking the individual private scores
- **Ring-Merging:** merging the small degree ciphertext of the individual scores into large degree ciphertexts
- **Scheme-Switching:** switching from coefficient packed messages to CKKS SIMD encoded messages
- **Private Thresholds with CKKS:** local and global private selection over the encrypted scores
- **Bootstrapping with CKKS** to evaluate the private thresholds

The implementation is available at

<https://github.com/Pro7ech/fhe-org-2024>

Let

- $f : \mathbb{R}_{[a,b]} \rightarrow \mathbb{R}$
- $g : \mathbb{R}_{[a,b]} \rightarrow \mathbb{R}_{[0,1]} : y = \frac{1}{2} \left( \frac{2x-b-a}{b-a} + 1 \right)$

Then, given a discretization factor of  $1/N$ , we can encode  $f$  on a polynomial of  $\mathbb{R}[X]/(X^N + 1)$  as

$$\mathcal{F}(X) = f'(0) - \sum_{i=1}^{N-1} f' \left( g^{-1} \left( \frac{N-i}{N} \right) \right) \cdot X^i$$

and for  $x \in [a, b]$ , we have

$$X^{\lfloor Ng(x) \rfloor} \cdot \text{RLWE}(\mathcal{F}(X)) \approx \text{RLWE}(f(x)X^0 + \star)$$

with an error bounded by  $|f(x) - f(x \pm \frac{b-a}{N})|$



This step is not a native operation in any of the HE schemes in LATTIGO .  
But It can be implemented from the lower-level packages `rlwe` and `ring`:

- 1 Encode  $f$  on a `[]float64` of size  $N$  with the method described above
- 2 Encode the `[]float64` on an `rlwe.Plaintext` using the `hefloat.Encoder`
- 3 Encrypt the `rlwe.Plaintext` on an `rlwe.Ciphertext` using the `rlwe.Encryptor`

The code for this step is located in the files

- `function.go`
- `server.go` (for the API call)



The evaluation step of the encrypted scoring function is done as:

$$\text{RLWE}\left(\sum_{j=0}^{h-1} f_j(x_j) X^0 + \star\right) \approx \sum_{j=0}^{h-1} \text{NTT}\left(X^{\lfloor Ng(x_j) \rfloor}\right) \cdot \text{RLWE}\left(\mathcal{F}_j(X)\right)$$

and can be carried out with basic polynomial arithmetic of the `ring` package:

- `ring.NTT(p0, p1 ring.Poly)`
- `ring.MulCoeffsMontgomery(p0, p1, p2 ring.Poly)`
- `ring.MulCoeffsMontgomeryThenAdd(p0, p1, p2 ring.Poly)`

- Each evaluation of **Encrypted Scoring Function** step returns an RLWE whose first coefficient is the evaluation of  $f(x_i)$
- We can homomorphically repack  $N$  of those coefficients in a single RLWE ciphertext:

$$\sum_{i=0}^{N-1} \text{RLWE}(m_i + \star) \cdot X^i \xrightarrow{\text{Repack}} \text{RLWE} \left( \sum_{i=0}^{N-1} m_i \cdot X^i \right)$$

- This reduces the number of ciphertexts by a factor of  $N$



LATTIGO has a native support of **Ring-Packing** through the API of the `rlwe.Evaluator` which enables to pack a hash map containing indexed `rlwe.Ciphertext` from 0 to  $N - 1$  into a single `rlwe.Ciphertext`:

- Method: `rlwe.Evaluator.Pack`
- Input: `map[int]*rlwe.Ciphertext`
- Output: `rlwe.Ciphertext`

The code for this step is located in the files

- `repacking.go`
- `repacking_keys.go`
- `server.go` (for the API call)

- The output of the **Ring-Packing** step returns RLWE ciphertexts whose ring degree is small, e.g.  $N = 2^{12}$
- **Scheme-Switching** to CKKS makes use of the CKKS bootstrapping, which requires a ring degree  $N' = 2^{16}$  (for security)
- We can merge  $N'/N$  RLWE ciphertexts of degree  $N$  into a single RLWE ciphertext of degree  $N'$ :

$$\sum_{j=0}^{N'/N-1} \text{RLWE} \left( \sum_{i=0}^{N-1} m_{jN+i} \cdot X^i \right) \cdot X^j$$

↓ Merge

$$\text{RLWE} \left( \sum_{j=0}^{N'/N-1} \left( \sum_{i=0}^{N-1} m_{jN+i} \cdot X^{iN'/N} \right) \cdot X^j \right)$$





LATTIGO has native support for switching RLWE ciphertext dimensions through key-switching and Ring-Merging can be implemented in only a few lines of code:

$$\text{RLWE}_{2N} \leftarrow \text{KeySwitch}_{\pi(s) \rightarrow s'}(\pi(\text{RLWE}_N) + X \cdot \pi(\text{RLWE}_N))$$

- $\pi : X \rightarrow X^2$  : `rlwe.SwitchCiphertextRingDegreeNTT`
- `KeySwitch` : `rlwe.Evaluator.ApplyEvaluationKey`

The code for this step is located in the files:

- `repacking.go`
- `repacking_keys.go`
- `server.go` (for the API call)



- The **Scheme Switching** step homomorphically encodes the values and raises the modulus to enable further computations
- This is equivalent to a CKKS bootstrapping that skips the last step (SlotsToCoeffs):
  - 1 **ModRaise**: raises the ciphertext to the largest modulus
  - 2 **CoeffsToSlots**: homomorphically encodes the ciphertext
  - 3 **EvalMod**: homomorphically evaluates the modular reduction
  - 4 (**SlotsToCoeffs**: homomorphically decodes the ciphertext)



LATTIGO supports for CKKS bootstrapping with out of the box default parameters, as well as advanced parameterization, enabling custom instantiation as well as custom circuit composition:

- `bootstrapping.Evaluator`
- `.ScaleDown`: scales the ciphertext down to  $Q_0$
- `.ModUp`: raises the ciphertext modulus from  $Q_0$  to  $Q_L$
- `.CoeffsToSlots`: homomorphic encoding
- `.EvalMod`: homomorphic modular reduction
- `.SlotsToCoeffs`: homomorphic decoding

The code for this step is located in the files

- `bootstrapping.go`
- `scheme_switching.go`
- `server.go` (for the API call)

Let

$$\text{Thresh}_t(x) : \begin{cases} 1 & \text{if } x \geq t \\ 0 & \text{otherwise} \end{cases} \quad \text{and} \quad \text{Sign}(x) : \begin{cases} 1 & \text{if } x > 0 \\ 0.5 & \text{if } x = 0 \\ -1 & \text{otherwise} \end{cases}$$

then  $\text{Thresh}_t(x) = \frac{1}{2}\text{Sign}(x - t) + \frac{1}{2}$  and we can evaluate Sign in the interval  $[-1, -2^{-\alpha}] \cup [2^{-\alpha}, 1]$  with precision  $\beta$  by using a composition of successive minimax approximations<sup>1</sup>

This approach requires to scales  $x - t$  to the interval  $[-1, 1]$ , which is done during the local threshold by a normalization factor  $1/\sum \max f_j$  and during the global threshold by a normalization factor  $1/p$

---

<sup>1</sup>Lee et al, *Minimax Approximation of Sign Function by Composite Polynomial for Homomorphic Comparison*

LATTIGO has native support for comparison, through the `hefloat.ComparisonEvaluator`. The user can generate and provide a custom mimimax composite approximation that suits the application needs

- **Local Private Threshold:**  $[-1, -2^{-8}] \cup [2^{-8}, 1]$  with precision  $\beta = 14$ , total depth of 12
- **Global Private Threshold:**  $[-1, -2^{-16}] \cup [2^{-16}, 1]$  with precision  $\beta = 9.4$ , total depth of 20
- `hefloat.ComparisonEvaluator.Step`

To avoid  $x = t_i$  in both cases, which would return 0.5, we add a small bias to  $x - t$  that is half the smallest difference between two possible values of  $x$

The code for this step is located in the files

- `server.go`

- 1 The client generates:
  - The necessary evaluation keys
  - The  $h$  encrypted scoring functions  $[\text{Enc}(f[0]), \dots, \text{Enc}(f[h-1])]$
  - The local encrypted threshold  $\text{Enc}(t_0)$  and normalization factor  $\text{Enc}(1/\sum \max(f_j))$
  - The global encrypted threshold  $\text{Enc}(t_1)$

and sends them to the server

- 2 The server evaluates  $\text{Enc}(y[i]) = \sum X^{\mathcal{P}[i][j]} \cdot \text{Enc}(f[j])$
- 3 The server packs the  $\text{Enc}(y[i])$  into RLWE low-degree ciphertexts
- 4 The server merge the low-degree RLWE ciphertexts into large-degree RLWE ciphertexts
- 5 The server scheme-switches the large-degree RLWE ciphertexts into CKKS ciphertexts
- 6 The server evaluates the private local-thresholds and aggregates all values
- 7 The server evaluates the final private global-threshold

# PERFORMANCE

- Set I: Encrypted Lookup-Table & Ring-Packing
- Set II-V: Ring-Merging
- Set V: Local & Global Thresholds
- Set VI: Scheme-Switching & Bootstrapping

Set	$\log(N)$	$\log(Q)$	$\log(P)$	$\log(QP)$	$h$	$\lambda$
I	12	60	48	108	$2N/3$	128
II	13	60	48	108	$2N/3$	$> 128$
III	14	60	48	108	$2N/3$	$> 128$
IV	15	60	48	108	$2N/3$	$> 128$
V	16	$60 + 8 \cdot 45$	158	578	192	$> 128$
VI	16	$60 + 8 \cdot 45 + 3 \cdot 39$ $+ 8 \cdot 60 + 4 \cdot 56$	$5 \cdot 61$	1546	192	128



Set	Size [MB]
Ring Packing Keys	3.25
Ring Merging Keys	7.5
Bootstrapping Keys	7395
$RLWE(f_j)$	2.01
$RLWE(t_0)$ & $RLWE(1/\sum \max(f_j))$	18
$RLWE(t_1)$	9
Total	$\approx 7435$

- Timings for a database of  $2^{19}$  entries single threaded<sup>2</sup>

Operation	Total	Amortized
Client Initialization	22.7s	43 $\mu$ s
Server Initialization	10.3s	19 $\mu$ s
Private Functions & Packing	436.6s	832 $\mu$ s
Scheme-Switching	64.2s	122 $\mu$ s
Private-Threshold Local	327.1s	624 $\mu$ s
Private-Threshold Global	49.9s	95 $\mu$ s
<b>Total</b>	<b>910.8s</b>	<b>1.737ms</b>

- Peak of 22GB of RAM, most of it being due to the size of the bootstrapping keys and the plaintext matrices for the homomorphic encoding/decoding during the bootstrapping

<sup>2</sup>i9-12900K, 64GB DDR5, Windows 11, Go 1.22

- In many real-world settings, data are distributed across several databases that may be independently managed. Computing the circuit without leaking becomes a secure multiparty computation (MPC) problem
- LATTIGO supports multiparty homomorphic encryption (MHE) schemes that enable an efficient MPC protocol. We have recently released **Helium**, a first implementation of such a protocol



<https://github.com/tuneinsight/lattigo>  
<https://github.com/ChristianMct/helium>